

Application 2: Analysis of a Computer Network

Algorithmic Thinking (Part1)

Shamsuddin Rehmani

July 16, 2019

Application 2 Description

Graph exploration (that is, “visiting” the nodes and edges of a graph) is a powerful and necessary tool to elucidate properties of graphs and quantify statistics on them. For example, by exploring a graph, we can compute its degree distribution, pairwise distances among nodes, its connected components, and centrality measures of its nodes and edges. As we saw in the Homework and Project, breadth-first search can be used to compute the connected components of a graph.

In this Application, we will analyze the connectivity of a computer network as it undergoes a cyber-attack. In particular, we will simulate an attack on this network in which an increasing number of servers are disabled. In computational terms, we will model the network by an undirected graph and repeatedly delete nodes from this graph. We will then measure the resilience of the graph in terms of the size of the largest remaining connected component as a function of the number of nodes deleted.

Answer to Question 1

Probability p such that the ER graph computed using this edge probability has approximately the same number of edges as the computer network is 0.00397

Integer m such that the number of edges in the UPA graph is close to the number of edges in the computer network is 2

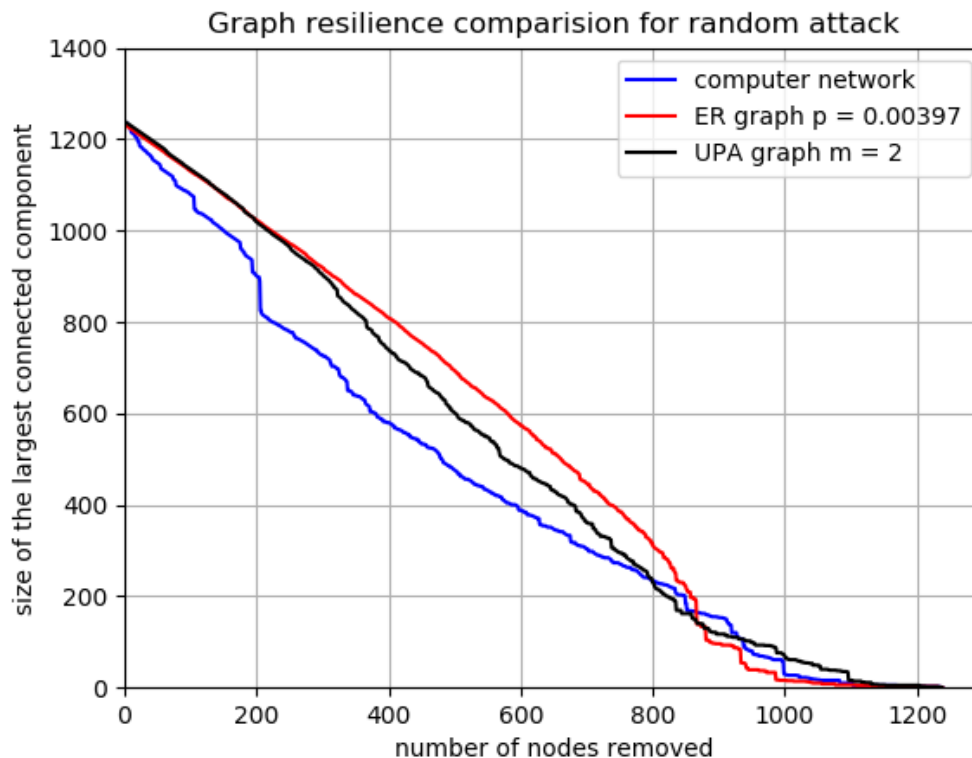


Figure 1: graph resilience comparison for random attack for computer network, undirected ER graph, and UPA graph

Answer to Question 2

All 3 graphs seem to be resilient (as shown in Figure 1), i.e the size of the largest connected component is within 25% of 1000 (the approximate number of remaining nodes)

Answer to Question 3

Big-O bounds for `target_order` = $O(n^2 + m) = O(n^2)$ since for UPA graph $m \leq 5n$ (m = total number of edges in UPA)

Big-O bounds for `fast_target_order` = $O(n + m) = O(n)$ since for UPA graph $m \leq 5n$ (m = total number of edges in UPA)




Figure 2: runtime for `target_order` and `fast_target_order` on UPA graph with nodes ranging from 10 to 1000 for $m = 5$

Answer to Question 4




Figure 3: graph resilience comparison for targeted attack for computer network, undirected ER graph, and UPA graph

Answer to Question 5

From the graph(Figure 3) we can see that only ER graph is resilient as the first 20% of the nodes are removed while the UPA and the computer network graph reaches close to zero as 20% of the nodes are removed

A Python code used to answer the Application Questions

```
1  """
2  Solutiion for Application #2: "Analysis of a Computer Network"
3  """
4
5  import time
6  import random
7  import matplotlib.pyplot as plt
8  import numpy as np
9  import alg_application2.provided as alg_app2_prov
10 import alg_example_graphs as alg_graphs
11 import alg_project2.solution as alg_proj2_sol
12
13 ##### Q1 Solution #####
14 # To begin our analysis, we will examine the resilience of the computer network under
15 # an attack in which servers are chosen at random. We will then compare the resilience
16 # of the network to the resilience of ER and UPA graphs of similar size.
17 #
18 # To begin, you should determine the probability pp such that the ER graph computed
19 # using this edge probability has approximately the same number of edges as the computer
20 # network. (Your choice for pp should be consistent with considering each edge in the
21 # undirected graph exactly once, not twice.) Likewise, you should compute an integer mm
22 # such that the number of edges in the UPA graph is close to the number of edges in the
23 # computer network. Remember that all three graphs being analyzed in this Application
24 # should have the same number of nodes and approximately the same number of edges.
25
26 # load the graph from a text file
27 cnet_graph = alg_app2_prov.load_graph(alg_app2_prov.NETWORK_URL)
28
29 # get the number of nodes in the computer network graph
30 num_nodes = len(cnet_graph.keys())
31
32 # find the total number of edges in th computer network graph
33 edges = sum([len(neighbors) for neighbors in cnet_graph.values()])/2
34
35 # find the probability such that the ER graph computed using this edge
36 # probability has approximately the same number of edges as the computer network
37 prob_p = round(2.0 * edges / (num_nodes * (num_nodes - 1.0)), 6)
38
39 # get the average degree so that the graph created using UPA algorithm has approximately
40 # same number of edges as network_graph
41 m_nodes = int(round(float(edges)/num_nodes))
42
43 # generate the random graph based on ER algorithm
44 er_graph = alg_graphs.alg_er(num_nodes, prob_p)
45
46 # generate the random graph based on UPA algorithm
47 upa_graph = alg_graphs.alg_upa(num_nodes, m_nodes)
48
49 # Next, you should write a function random_order that takes a graph and returns a list
50 # of the nodes in the graph in some random order. Then, for each of the three graphs
51 # (computer network, ER, UPA), compute a random attack order using random_order and use
52 # this attack order in compute_resilience to compute the resilience of the graph.
53
54 def random_order(graph):
55     """
56     Take a graph a returns a random sequence of its nodes
57     Arguments:
58         graph {dictionary} — [a graph]
59
60     Returns:
61         list of nodes — random sequence of nodes
62     """
63
64     lst_nodes = graph.keys()
65     random.shuffle(lst_nodes)
66
67     return lst_nodes
68
69 # compute the resilience of each of the 3 graphs
70 cnet_res = alg_proj2_sol.compute_resilience(cnet_graph, random_order(cnet_graph))
71 er_res = alg_proj2_sol.compute_resilience(er_graph, random_order(er_graph))
72 upa_res = alg_proj2_sol.compute_resilience(upa_graph, random_order(upa_graph))
73
74 # Once you have computed the resilience for all three graphs, plot the results as three
75 # curves combined in a single standard plot (not log/log). Use a line plot for each curve.
76 # The horizontal axis for your single plot be the number of nodes removed (ranging
77 # from zero to the number of nodes in the graph) while the vertical axis should be the
78 # size of the largest connect component in the graphs resulting from the node removal.
79 # For this question (and others) involving multiple curves in a single plot, please
80 # include a legend in your plot that distinguishes the three curves. The text labels in
81 # this legend should include the values for pp and mm that you used in computing the ER
82 # and UPA graphs, respectively. Both matplotlib and simpleplot support these capabilities
83 # (matplotlib example and simpleplot example).
84 #
85 # Note that three graphs in this problem are large enough that using CodeSkulptor to
86 # calculate compute_resilience for these graphs will take on the order of 3–5 minutes
87 # per graph. When using CodeSkulptor, we suggest that you compute resilience for each
88 # graph separately and save the results (or use desktop Python for this part of the
89 # computation). You can then plot the result of all three calculations using simpleplot.
90 # load the graph from the text file
91
92 # compute the list of number of nodes removed (ranging from zero to the number of nodes in the graph)
93 num_removed = range(num_nodes + 1)
94
```

```

95 # plot the graphs of resilience vs number of nodes removed for each of the 3 graphs
96 plt.figure(0)
97 plt.plot(num_removed, cnet_res, '-b', label = 'computer network')
98 plt.plot(num_removed, er_res, '-r', label = 'ER graph p = 0.00397')
99 plt.plot(num_removed, upa_res, '-k', label = 'UPA graph m = 2')
100 plt.title('Graph resilience comparison for random attack')
101 plt.xlabel('number of nodes removed')
102 plt.ylabel('size of the largest connected component')
103 plt.legend(loc = 'upper right')
104 plt.xlim(0, None)
105 plt.ylim(0, 1400)
106 plt.grid()
107 # uncomment to save the plot
108 #plt.savefig("Q1_graph_resilience_comparison.png")
109
110 ##### Q2 Solution #####
111 # Consider removing a significant fraction of the nodes in each graph
112 # using random_order. We will say that a graph is resilient under this
113 # type of attack if the size of its largest connected component is
114 # roughly (within ~25%) equal to the number of nodes remaining, after
115 # the removal of each node during the attack.
116 #
117 # Examine the shape of the three curves from your plot in Question 1.
118 # Which of the three graphs are resilient under random attacks as the
119 # first 20% of their nodes are removed?
120 #
121 # Ans: all 3 graphs seem to be resilient, i.e the size of the largest
122 # connected component is within 25% of 1000 ( the approximate number of
123 # remaining nodes)
124
125 ##### Q3 Solution #####
126 # In the next three problems, we will consider attack orders in which the
127 # nodes being removed are chosen based on the structure of the graph. A
128 # simple rule for these targeted attacks is to always remove a node of
129 # maximum (highest) degree from the graph. The function targeted_order(ugraph)
130 # in the provided code takes an undirected graph ugraph and iteratively does
131 # the following:
132 #
133 # - Computes a node of the maximum degree in ugraph. If multiple nodes have
134 #   the maximum degree, it chooses any of them (arbitrarily).
135 # - Removes that node (and its incident edges) from ugraph.
136 #
137 # Observe that targeted_order continuously updates ugraph and always computes
138 # a node of maximum degree with respect to this updated graph. The output of
139 # targeted_order is a sequence of nodes that can be used as input to compute_resilience.
140 #
141 # As you examine the code for targeted_order, you feel that the provided
142 # implementation of targeted_order is not as efficient as possible. In
143 # particular, much work is being repeated during the location of nodes
144 # with the maximum degree. In this question, we will consider an alternative
145 # method (which we will refer to as fast_targeted_order) for computing the
146 # same targeted attack order. In Python, this method creates a list
147 # degree_sets whose kth element is the set of nodes of degree k. The method
148 # then iterates through the list degree_sets in order of decreasing degree.
149 # When it encounters a non-empty set, the nodes in this set must be of
150 # maximum degree. The method then repeatedly chooses a node from this set,
151 # deletes that node from the graph, and updates degree_sets appropriately.
152 #
153 # For this question, your task is to implement fast_targeted_order and then
154 # analyze the running time of these two methods on UPA graphs of size n with
155 # m = 5.
156
157 # Determine big-O bounds of the worst-case running times of targeted_order
158 # and fast_targeted_order as a function of the number of nodes n in the UPA graph.
159 # Since the number of edges in these UPA graphs is always less than 5n (due to the
160 # choice of m = 5), your big-O bounds for both functions should be expressions in n.
161 # You should also assume that all of the set operations used in fast_targeted_order
162 # are O(1).
163 #
164 # Ans: target_order = O(n^2 + m) = O(n^2) since for UPA graph m ≤ 5n (m = total
165 #       number of edges in UPA)
166 #       fast_target_order = O(n + m) = O(n) since for UPA graph m ≤ 5n (m = total
167 #       number of edges in UPA)
168
169 # Next, run these two functions on a sequence of UPA graphs with n in range(10,1000,10)
170 # and m=5 and use the time module (or your favorite Python timing utility) to compute
171 # the running times of these functions. Then, plot these running times (vertical axis)
172 # as a function of the number of nodes n (horizontal axis) using a standard plot
173 # (not log/log). Your plot should consist of two curves showing the results of your
174 # timings. Remember to format your plot appropriately and include a legend. The title
175 # of your plot should indicate the implementation of Python (desktop Python vs. CodeSkulptor)
176 # used to generate the timing results.
177
178 def fast_targeted_order(graph):
179     """
180     Compute a targeted attack order consisting of nodes of
181     maximal degree. The algorithm used was provided.
182
183     Arguments:
184         graph {dictionary} — a graph
185
186     Returns:
187         list of nodes — a list of nodes of attack order with maximal
188         degree in descending order
189     """
190

```

```

191     # initialize the list with the target node order
192     node_order = []
193
194     # make a copy of the node
195     graph_cpy = alg_app2_prov.copy_graph(graph)
196
197     # get all the nodes of the graph
198     nodes = graph_cpy.keys()
199
200     # initialize all degree set so that all degree corresponds to empty set of nodes
201     degree_set = [set() for dummy_idx in nodes]
202
203     # add all the nodes to their corresponding degree location
204     for node in nodes:
205         degree = len(graph_cpy[node])
206         degree_set[degree].add(node)
207
208     # update the degree set and delete the node of the copy of graph appropriately after
209     # storing the node with current maximal degree
210     for deg_set_idx in range(len(nodes) - 1, -1, -1):
211         while(len(degree_set[deg_set_idx]) != 0):
212             node_u = degree_set[deg_set_idx].pop()
213             for neighbor in graph_cpy[node_u]:
214                 degree_neighbor = len(graph_cpy[neighbor])
215                 degree_set[degree_neighbor].remove(neighbor)
216                 degree_set[degree_neighbor - 1].add(neighbor)
217
218             node_order.append(node_u)
219             alg_app2_prov.delete_node(graph_cpy, node_u)
220
221     return node_order
222
223 # initialize the fast and normal function times for targeted order
224 time_fast_targeted_order = []
225 time_targeted_order = []
226
227 # initialize the nodes to be used to generate the UPA graphs
228 nodes = range(10, 1000, 10)
229
230 # calculate the time to run the normal and fast functions for the target order
231 for node in nodes:
232     # create the UPA graph with n = node and m = 5 where m is the number of
233     # existing nodes to which a new node is connected during each iteration
234     upa_graph_new = alg_graphs.alg_upa(node, 5)
235     # calculate the attack order based on normal targeted order function
236     # and store the time it takes to run this function
237     start = time.time()
238     alg_app2_prov.targeted_order(upa_graph_new)
239     end = time.time()
240     time_targeted_order.append((end - start) * 1000)
241     # calculate the attack order based on fast targeted order function
242     # and store the time it takes to run this function
243     start = time.time()
244     fast_targeted_order(upa_graph_new)
245     end = time.time()
246     time_fast_targeted_order.append((end - start) * 1000)
247
248 # plot the graphs of resilience vs number of nodes removed for each of the 3 graphs
249 plt.figure(1)
250 plt.plot(nodes, time_targeted_order, '-b', label = 'targeted_order')
251 plt.plot(nodes, time_fast_targeted_order, '-k', label = 'fast_targeted_order')
252 plt.title('regular vs fast runtime of targeted order in Visual Studio')
253 plt.xlabel('number of nodes of UPA graph for m = 5')
254 plt.ylabel('run times[msec]')
255 plt.legend(loc = 'upper left')
256 plt.xlim(10, None)
257 plt.ylim(0, None)
258 plt.grid()
259
260 # uncomment to save the plot
261 #plt.savefig("Q3-targeted_order_time_comparision.png")
262
263 ##### Q4 Solution #####
264 # To continue our analysis of the computer network, we will examine its resilience
265 # under an attack in which servers are chosen based on their connectivity. We will
266 # again compare the resilience of the network to the resilience of ER and UPA graphs
267 # of similar size.
268 #
269 # Using targeted_order (or fast_targeted_order), your task is to compute a targeted
270 # attack order for each of the three graphs (computer network, ER, UPA) from Question 1.
271 # Then, for each of these three graphs, compute the resilience of the graph using
272 # compute_resilience. Finally, plot the computed resilience as three curves (line plots)
273 # in a single standard plot. As in Question 1, please include a legend in your plot that
274 # distinguishes the three plots. The text labels in this legend should include the values
275 # for p and m that you used in computing the ER and UPA graphs, respectively.
276
277 # compute the target order for the 3 graph in Q1
278 tar_order_cnet = fast_targeted_order(cnet_graph)
279 tar_order_er = fast_targeted_order(er_graph)
280 tar_order_upa = fast_targeted_order(upa_graph)
281
282 # compute the resilience for the 3 graph using the targeted order
283 # compute the resilience of each of the 3 graphs
284 cnet_res = alg_proj2_sol.compute_resilience(cnet_graph, tar_order_cnet)
285 er_res = alg_proj2_sol.compute_resilience(er_graph, tar_order_er)
286 upa_res = alg_proj2_sol.compute_resilience(upa_graph, tar_order_upa)

```

```

287
288 # plot the computer resilience for the targeted order for the 3 graph
289 plt.figure(2)
290 plt.plot(num_removed, cnet_res, '-b', label = 'computer network')
291 plt.plot(num_removed, er_res, '-r', label = 'ER graph p = 0.00397')
292 plt.plot(num_removed, upa_res, '-k', label = 'UPA graph m = 2')
293 plt.title('Graph resilience comparision for targeted order')
294 plt.xlabel('number of nodes removed')
295 plt.ylabel('size of the largest connected component')
296 plt.legend(loc = 'upper right')
297 plt.xlim(0, None)
298 plt.ylim(0, 1400)
299 plt.grid()
300 plt.show()
301 # uncommet to save the plot
302 #plt.savefig("Q4-graph-resilience-comparision.png")
303
304 ##### Q5 Solution #####
305 # Examine the shape of the three curves from your plot in Question 4.
306 # Which of the three graphs are resilient under targeted attacks as
307 # the first 20% of their nodes are removed? Again, note that there is
308 # no need to compare the three curves against each other in your answer
309 # to this question.
310 #
311 # Ans: From the graph we can see that only ER graph is resilient as the
312 # first 20% of the nodes are removed while the UPA and the computer
313 # network graph reaches close to zero as 20% of the ndoes are removed

```

B All functions for project 4 used in the application

```

1 """
2 Functions for Project #2: "Connected Components and Graph Resilience". These
3 functions will be used for the Application #2: "Analysis of a Computer Network"
4 """
5 from collections import deque
6 import alg_application2_provided as alg_app2_prov
7
8 def bfs_visited(ugraph, start_node):
9     """
10     Performs the Breath First Search(BFS) and returns a set of all the nodes
11     that are visited starting from start_node
12
13     Arguments:
14         ugraph {dictionary} — an undirect graph
15         start_node {integer} — a node in the ugraph
16
17     Returns:
18         set — a set of all nodes visited by a BFS that start at start_node
19     """
20     # initialize an queue with the start_node. We use python's built in double
21     # ended queue, deque
22     deq = deque([start_node])
23     # add the start node to the visited nodes set
24     visited = set([start_node])
25     # keep traversing through all the neighbors of the nodes in the queue
26     # as long as the queue is not empty and mark them as visited if the nodes
27     # are not yet visited
28     while len(deq) != 0:
29         curr_node = deq.popleft()
30         for neighbor in ugraph[curr_node]:
31             if not (neighbor in visited):
32                 visited.add(neighbor)
33                 deq.append(neighbor)
34
35     return visited
36
37 def cc_visited(ugraph):
38     """
39     Takes an undirected graph ugraph and computes the all the
40     connected components of the graph
41
42     Arguments:
43         ugraph {dictionary} — an undirected graph
44
45     Returns:
46         list of sets — returns a list of sets where each set has all the nodes in
47                         a particular connected component of the graph, and each set
48                         represent a connect component of the graph
49     """
50     # initialize the remaining nodes in the ugraph that have not yet been visited
51     remain_nodes = set(ugraph.keys())
52     # initiakize the list of sets where each set is a connected component of ugraph
53     con_comp = []
54
55     # use BFS to find all the connect components until all the nodes of the ugraph
56     # have been visisted
57     while len(remain_nodes) != 0:
58         not_vis_node = remain_nodes.pop()
59         visited = bfs_visited(ugraph, not_vis_node)
60         con_comp.append(visited)
61         remain_nodes -= visited
62
63     return con_comp
64

```

```

65 def largest_cc_size(ugraph):
66     """
67     Takes a undirected graph and returns the size of the largest connected component
68
69     Arguments:
70         ugraph {dictionary} — an undirected graph
71
72     Returns:
73         integer — the size of the largest connected component of ugraph
74     """
75     # find the size of all the connect components of the ugraph
76     len_cc = [len(con_comp) for con_comp in cc_visited(ugraph)]
77
78     # make sure to take care of the case when ugraph is empty and we get
79     # an empty len_cc list
80     if (len(len_cc) == 0):
81         return 0
82
83     # return the max size of connected compo
84     return max(len_cc)
85
86 def compute_resilience(ugraph, attack_order):
87     """
88     Computes a measure of resilience of an undirected graph. Takes the undirected
89     graph ugraph, a list of nodes attack_order and iterates through the nodes in
90     attack_order. For each node in the list, the function removes the given node
91     and its edges from the graph and then computes the size of the largest connected
92     component for the resulting graph.
93
94     Arguments:
95         ugraph {dictionary} — an undirected graph
96         attack_order {list of nodes} — list of nodes that will be iterated over
97
98     Returns:
99         list of integers — return a list whose (k+1)th entry is the size of the largest
100                           connected component in the graph after the removal of the first
101                           k nodes in attack_order
102     """
103     new_graph = alg_app2_prov.copy_graph(ugraph)
104
105     # get the size of the largest connected component before removing any nodes
106     lst_max_cc = [largest_cc_size(new_graph)]
107
108     # start removing each node in the attack_order and its edges from the ugraph
109     # and find the largest connected component after each removal
110     for remove_node in attack_order:
111         alg_app2_prov.delete_node(new_graph, remove_node)
112         lst_max_cc.append(largest_cc_size(new_graph))
113
114     return lst_max_cc

```

C Code for ER and UPA graph generations

```

1  """
2  Funtions to generate 2 types of ugraphs, the undirected ER graph
3  and UPA graph
4  """
5  import random
6  import alg_upa_trial as upa
7
8  def make_complete_graph(num_nodes):
9      """
10      create and return a complete graph with nodes from
11      0 to num_nodes - 1 for num_nodes > 0. Otherwise
12      the function returns a dictionary corresponding to
13      the empty graph
14
15      Arguments:
16          num_nodes {integer} — number of nodes for the graph
17
18      Returns:
19          dictionary — returns a dictionary corresponding to a complete directed
20                      graph with the specified number of nodes.
21      """
22      # local variable for the complete graph
23      graph = {}
24
25      # return an empty graph if num_nodes is not positive
26      if num_nodes <= 0:
27          return graph
28
29      for node in range(num_nodes):
30          # create an adjacency list for a directed complete graph with no
31          # self loops or parallel edges
32          graph[node] = set([val for val in range(num_nodes) if val != node])
33
34      return graph
35
36
37 def alg_er(num_nodes, p):
38     """
39     generate a random graph based on Erdos Renyi(ER) model G(n, p)
40     where each edge in the graph is added with probability p
41

```



```

42 Arguments:
43     num_nodes {integer} — the total number of nodes for the generated graph
44     p {float} — the probability with which to add each edge to the generated graph
45
46 Returns:
47     dictionary — return the ER random graph
48
49
50 ugraph = {}
51 all_edges = []
52
53 if (num_nodes <= 0):
54     return ugraph
55
56 # create a graph of all nodes but no edges
57 for node in range(num_nodes):
58     ugraph[node] = set()
59
60 # find all possible edges of the graph
61 all_edges = (set([frozenset([node1, node2]) for node1 in range(num_nodes)
62                     for node2 in range(num_nodes) if node1 != node2]))
63
64 # convert each edge from frozenset to list for indexing
65 all_edges = [list(item) for item in all_edges]
66
67 # add edges to the graph with probability p
68 for edge in all_edges:
69     rand_prob = random.uniform(0,1)
70     if rand_prob < p:
71         ugraph[edge[0]].add(edge[1])
72         ugraph[edge[1]].add(edge[0])
73
74 return ugraph
75
76 def alg_upa(n_nodes, m_nodes):
77     """
78     Uses the DPA algorithm provided in Q3 of the Application #1
79     to generates a random undirected graph iteratively, where
80     each iteration a new node is created, added to the graph,
81     and connected to the subset of the existing node
82
83     Arguments:
84         n_nodes {integer} — final number of nodes in the generated graph
85         m_nodes {integer} — number of existing nodes to which a new node is connected
86                             during each iteration
87
88     Returns:
89         dictionary — the generated graph based on modified DPA algorithm for undirected graph
90
91     """
92     # create a complete graph of m_nodes nodes
93     upa_graph = make_complete_graph(m_nodes)
94
95     # create the UPATrial object corresponding to complete graph
96     upa_trial = upa.UPATrial(m_nodes)
97
98     # add each new node to the existing graph randomly
99     # chosen with probability:
100     # (in-degree of new_node + 1) / (in-degree of all nodes +
101     # total number of existing nodes)
102     # simulated by the run_trial of the UPATrial class
103     for new_node in range(m_nodes, n_nodes):
104         # randomly select m_nodes from the existing graph that
105         # the new_node will be connected to. Remove if any
106         # duplicate nodes in the m_nodes selected
107         neighbors = upa_trial.run_trial(m_nodes)
108
109         # update the existing graph to add this new node and its
110         # neighbors
111         upa_graph[new_node] = neighbors
112
113         # add this new node to all the neighbor nodes since this
114         # is a undirected graph
115         for neighbor in neighbors:
116             upa_graph[neighbor].add(new_node)
117
118     return upa_graph

```